



Same convex loss, three optimisers reach the same optimum

L-BFGS (superlinear)



IRLS / Newton (quadratic)



Gradient Descent (linear)

